

# Tsunami Prediction - MDS Datathon (Winner)

This notebook presents my winning solution for the **2025 UBC MDS Tsunami Prediction Datathon**.

The competition objective was to build a machine learning model that predicts whether an earthquake would trigger a tsunami based on seismic and geographic features.

## Executive Summary

This project combines machine learning with geospatial feature engineering to improve tsunami prediction accuracy

## Key Contributions

- Engineered geo-spatial features from raw latitude and longitude, including:
  - Distance to coastline (via Mercator-projected nearest-boundary calculation)
  - A classification of whether an earthquake happens under the land or sea
  - A tectonic plate each earthquake belongs to using PB2002 plate polygons
- Visualized data points along with earth map and tectonic plates to extract meaningful spatial insights
- Applied a structured machine learning workflow including:
  - Column Transformer: Scaling & OneHotEncoder
  - Cross-validated model comparison across 7 different ML algorithms
  - Hyperparameter tuning with RandomizedCV
- Selected a tuned Random Forest model achieving 0.87 ROC-AUC on CV

## Why This Solution Won

Although other teams achieved slightly higher raw accuracy scores, this solution stood out for:

- Strong understanding of earthquakes and how to predict tsunami leveraging geosciences domain knowledge
- Strong feature engineering
- Clear explanation of why raw latitude/longitude can cause overfitting
- Clean geospatial transformations grounded in real-world tsunami dynamics
- Transparent and interpretable modeling steps

This notebook has been cleaned and expanded after the competition to improve clarity and reproducibility. The competition submission itself was created under a 3-hour time

limit.

## Import Datasets & Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import altair_ally as aly

from sklearn.dummy import DummyClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import (
    GridSearchCV,
    RandomizedSearchCV,
    cross_validate,
    train_test_split,
)
from sklearn.naive_bayes import BernoulliNB, MultinomialNB, GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from scipy.stats import loguniform, randint
from sklearn.compose import make_column_transformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler

# For Lat and Long Data Preprocessing
import geopandas as gpd
from shapely.geometry import Point
import geodatasets
import cartopy.crs as ccrs
import cartopy.feature as cfeature
```

```
In [2]: # Load the datasets
test_df = pd.read_csv("../data/test.csv")
train_df = pd.read_csv("../data/train.csv")
```

## Data Visualization

```
In [3]: # Codes for visualizing data points along with plate boundaries on the map
# colored by a numeric feature
def plot_map_numeric(df, value_col, lon_col="longitude", lat_col="latitude",
                     cmap="viridis", markersize=6, title=None):
    """
    Plot latitude/longitude points on a world map colored by a NUMERICAL col
    """
    plate_boundaries_url = "https://raw.githubusercontent.com/fraxen/tectonics/master/plates.gpkg"
    plates = gpd.read_file(plate_boundaries_url)
```

```

plates = plates.to_crs("EPSG:4326")

fig = plt.figure(figsize=(14, 10))
ax = plt.axes(projection=ccrs.PlateCarree())

# Background
ax.add_feature(cfeature.LAND, facecolor="lightgray")
ax.add_feature(cfeature.OCEAN, facecolor="lightblue")
ax.add_feature(cfeature.COASTLINE)
ax.add_feature(cfeature.BORDERS, linestyle=":")

# Plot plate boundaries
plates.plot(ax=ax, color="grey", linewidth=1, transform=ccrs.PlateCarree)

# Scatter with continuous colormap
sc = ax.scatter(
    df[lon_col],
    df[lat_col],
    c=df[value_col],
    cmap=cmap,
    s=markersize,
    transform=ccrs.PlateCarree()
)

# Colorbar
cbar = plt.colorbar(sc, shrink=0.6)
cbar.set_label(value_col, fontsize=12)

# Title
if title is None:
    title = f"Map Colored by Numerical Variable '{value_col}'"
plt.title(title, fontsize=16)

plt.show()

```

In [4]:

```

# Codes for visualizing data points along with plate boundaries on the map
# colored by a categorical feature
def plot_map_categorical(df, category_col, lon_col="longitude", lat_col="latitude",
                        cmap="tab20", markersize=6, title=None):
    """
    Plot latitude/longitude points on a world map colored by a CATEGORICAL column.
    Automatically assigns colors using factorize().
    """
    plate_boundaries_url = "https://raw.githubusercontent.com/fraxen/tectonics/master/plate_boundaries.shp"
    plates = gpd.read_file(plate_boundaries_url)
    plates = plates.to_crs("EPSG:4326")

    # Convert categories → integer IDs
    cat_codes, uniques = pd.factorize(df[category_col])

    fig = plt.figure(figsize=(14, 10))
    ax = plt.axes(projection=ccrs.PlateCarree())

    # Background
    ax.add_feature(cfeature.LAND, facecolor="lightgray")
    ax.add_feature(cfeature.OCEAN, facecolor="lightblue")

```

```

    ax.add_feature(cfeature.COASTLINE)
    ax.add_feature(cfeature.BORDERS, linestyle=":")

    # Plot plate boundaries
    plates.plot(ax=ax, color="grey", linewidth=1, transform=ccrs.PlateCarree())

    # Scatter with discrete colormap
    sc = ax.scatter(
        df[lon_col],
        df[lat_col],
        c=cat_codes,
        cmap=cmap,
        s=markersize,
        transform=ccrs.PlateCarree()
    )

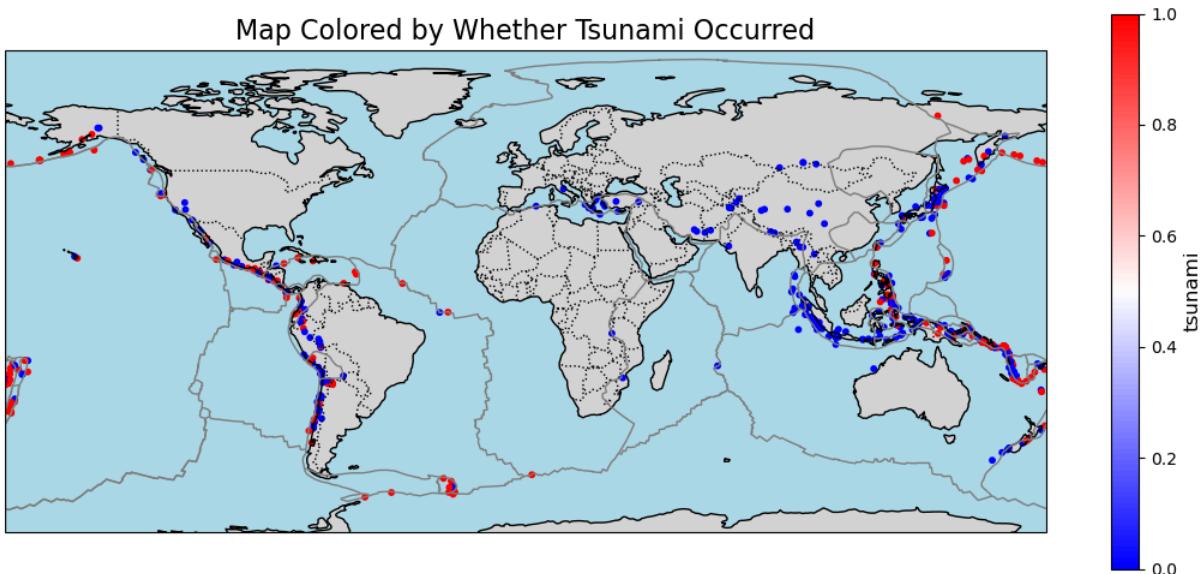
    # Categorical colorbar labels
    cbar = plt.colorbar(sc, shrink=0.6, ticks=range(len(uniques)))
    cbar.ax.set_yticklabels(uniques)
    cbar.set_label(category_col, fontsize=12)

    # Title
    if title is None:
        title = f"Map Colored by Categorical Variable '{category_col}'"
    plt.title(title, fontsize=16)

    plt.show()

```

In [27]: `plot_map_numeric(train_df, "tsunami", lon_col="longitude", lat_col="latitude", cmap="bwr", markersize=10, title="Map Colored by Whether Tsunami Occurred")`



## Transform Latitude and Longitude Columns

If we use latitude and longitude directly, it can lead the model to overfit. For a tsunami to happen, it is likely to be associated with whether an earthquake happens under the sea,

how far away it is from the coast, and the tectonic plate where an earthquake happens. Therefore, these transformations need to be done to account for that.

```
In [6]: # Load Natural Earth low-res land polygons via geodatasets
world_path = geodatasets.get_path("naturalearth.land")
land = gpd.read_file(world_path)

# Convert your lat/lon into point geometry
train_df['geometry'] = train_df.apply(lambda row: Point(row['longitude'], row['latitude']), axis=1)

# Make GeoDataFrame
gdf = gpd.GeoDataFrame(train_df, geometry='geometry', crs="EPSG:4326")

# Spatial join: which points intersect land?
joined = gpd.sjoin(gdf, land, how='left', predicate='intersects')

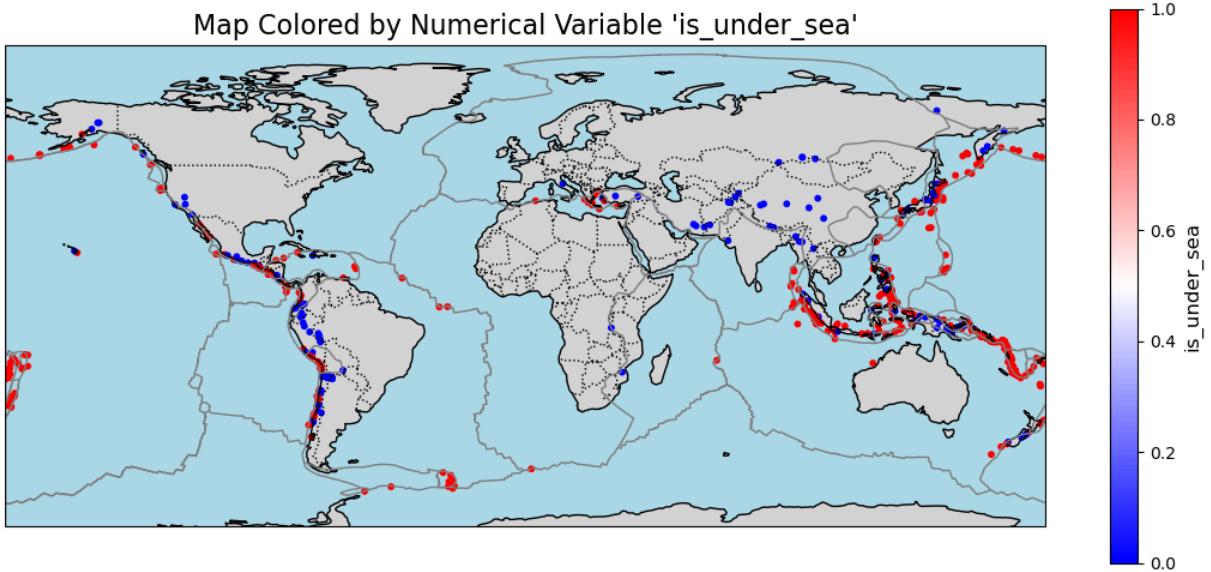
# If land data missing → ocean
train_df['is_under_sea'] = joined['featurecla'].isna().astype(int)

# Drop geometry if not needed
train_df = train_df.drop(columns=['geometry'])
train_df
```

```
Out[6]:   magnitude  cdi  mmi  sig  nst  dmin  gap  depth  latitude  longitude  Year
0          6.9    0     6   732   117  0.000  0.0   28.00 -4.0220  101.7760  2001
1          6.8    0     5   711    64  0.000  0.0   33.00  6.6310  126.8990  2001
2          7.2    4     4   838   443  0.000  17.0   19.00  2.4330  93.2100  2012
3          6.5    0     6   650   700  0.000  47.4   43.00  51.1480  157.5220  2006
4          6.6    8     8   831   782  0.000  16.7   12.00  37.5350  138.4460  2007
...
542         6.6    9     8  2840      0  0.174  25.0    8.00  42.8621  13.0961  2016
543         7.0    4     7   768   537  0.000  23.1   22.00 -2.1300  99.6270  2007
544         7.8    9     9  1545      0  0.481  21.0   15.11 -42.7373  173.0540  2016
545         6.5    7     7   756   178  0.430  54.0   10.00  23.0290  121.3480  2022
546         6.5    8     7  1400   578  0.000  38.6   59.00  17.9860 -99.7890  2011
```

547 rows × 14 columns

```
In [7]: plot_map_numeric(train_df, "is_under_sea", lon_col="longitude", lat_col="latitude", cmap="bwr", markersize=10, title=None)
```



```
In [8]: # Load Natural Earth land polygons
land_path = geodatasets.get_path("naturalearth.land")
land = gpd.read_file(land_path)

# Convert to coastline by taking the boundary
coast = land.boundary

# Project to world mercator (meters)
coast = coast.to_crs(epsg=3395)

def compute_distance_to_coast(row):
    # lon, lat → point
    point = Point(row["longitude"], row["latitude"])

    # project point to mercator
    point_m = gpd.GeoSeries([point], crs="EPSG:4326").to_crs(epsg=3395)[0]

    # shortest coastline distance
    dist_m = coast.distance(point_m).min()

    return dist_m / 1000 # convert to km

train_df["distance_to_coast_km"] = train_df.apply(compute_distance_to_coast,
train_df
```

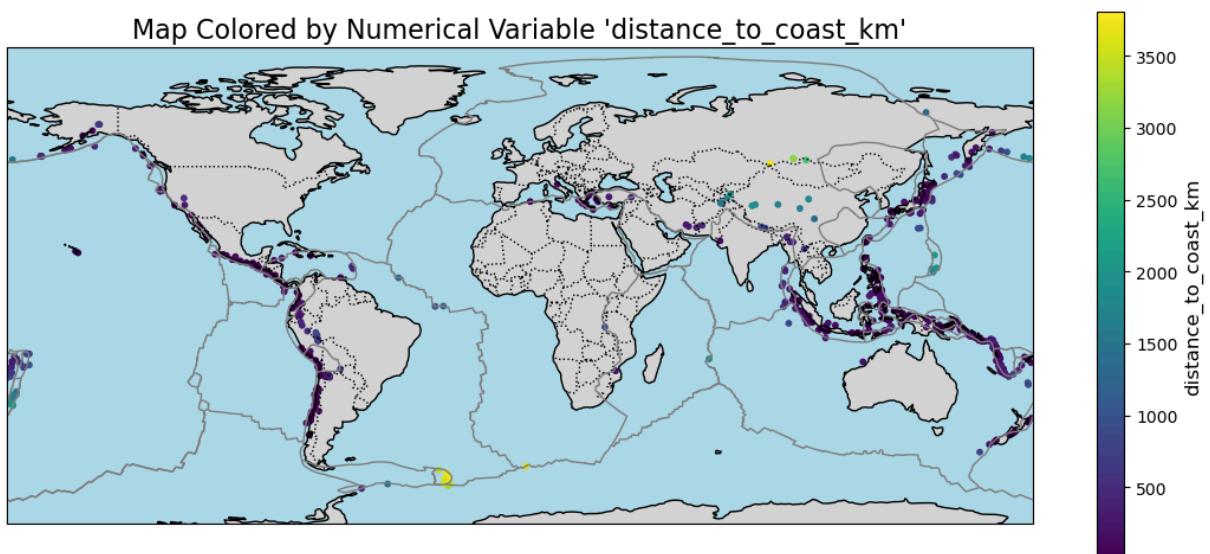
Out[8]:

	magnitude	cdi	mmi	sig	nst	dmin	gap	depth	latitude	longitude	Year
0	6.9	0	6	732	117	0.000	0.0	28.00	-4.0220	101.7760	2001
1	6.8	0	5	711	64	0.000	0.0	33.00	6.6310	126.8990	2001
2	7.2	4	4	838	443	0.000	17.0	19.00	2.4330	93.2100	2012
3	6.5	0	6	650	700	0.000	47.4	43.00	51.1480	157.5220	2006
4	6.6	8	8	831	782	0.000	16.7	12.00	37.5350	138.4460	2007
...	...	...	...	...	...	...	...	...	...	...	...
542	6.6	9	8	2840	0	0.174	25.0	8.00	42.8621	13.0961	2016
543	7.0	4	7	768	537	0.000	23.1	22.00	-2.1300	99.6270	2007
544	7.8	9	9	1545	0	0.481	21.0	15.11	-42.7373	173.0540	2016
545	6.5	7	7	756	178	0.430	54.0	10.00	23.0290	121.3480	2022
546	6.5	8	7	1400	578	0.000	38.6	59.00	17.9860	-99.7890	2011

547 rows × 15 columns

In [9]:

```
plot_map_numeric(train_df, "distance_to_coast_km", lon_col="longitude", lat_col="latitude", cmap="viridis", markersize=10, title=None)
```



In [10]:

```
plate_polygons_url = "https://raw.githubusercontent.com/fraxen/tectonicplates/master/GeoJSON/PB2015_0.5m.tectonicplates.json"
plates = gpd.read_file(plate_polygons_url)

plates = plates.to_crs("EPSG:4326") # ensure lat/long CRS

gdf = gpd.GeoDataFrame(
    train_df,
    geometry=gpd.points_from_xy(train_df["longitude"], train_df["latitude"]),
    crs="EPSG:4326"
)
```

```

train_df = gpd.sjoin(gdf, plates, predicate="within", how="left")
train_df = train_df.drop(columns=['geometry', 'index_right', 'LAYER', 'Code'])
train_df

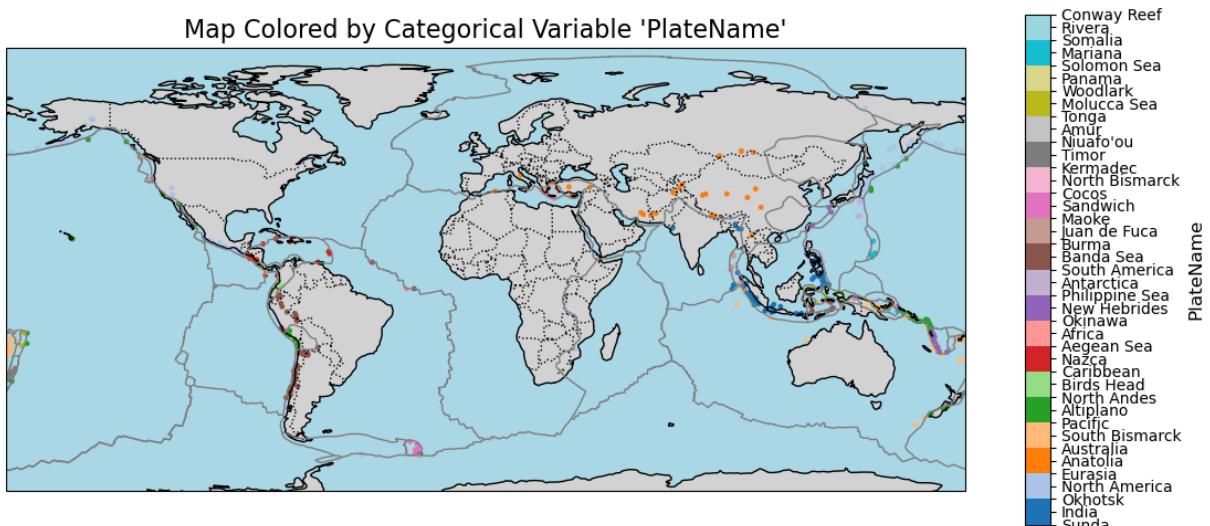
```

Out[10]:

	<b>magnitude</b>	<b>cdi</b>	<b>mmi</b>	<b>sig</b>	<b>nst</b>	<b>dmin</b>	<b>gap</b>	<b>depth</b>	<b>latitude</b>	<b>longitude</b>	<b>Year</b>
0	6.9	0	6	732	117	0.000	0.0	28.00	-4.0220	101.7760	2001
1	6.8	0	5	711	64	0.000	0.0	33.00	6.6310	126.8990	2001
2	7.2	4	4	838	443	0.000	17.0	19.00	2.4330	93.2100	2012
3	6.5	0	6	650	700	0.000	47.4	43.00	51.1480	157.5220	2006
4	6.6	8	8	831	782	0.000	16.7	12.00	37.5350	138.4460	2007
...	...	...	...	...	...	...	...	...	...	...	...
542	6.6	9	8	2840	0	0.174	25.0	8.00	42.8621	13.0961	2016
543	7.0	4	7	768	537	0.000	23.1	22.00	-2.1300	99.6270	2007
544	7.8	9	9	1545	0	0.481	21.0	15.11	-42.7373	173.0540	2016
545	6.5	7	7	756	178	0.430	54.0	10.00	23.0290	121.3480	2022
546	6.5	8	7	1400	578	0.000	38.6	59.00	17.9860	-99.7890	2011

547 rows × 16 columns

In [11]: `plot_map_categorical(train_df, "PlateName", lon_col="longitude", lat_col="latitude", cmap="tab20", markersize=6, title=None)`



## EDA

In [12]: `# The 12 original numerical predictor features`  
`numerical_features = [`  
`'magnitude', 'cdi', 'mmi', 'sig', 'nst', 'dmin', 'gap', 'depth', 'is_unc`

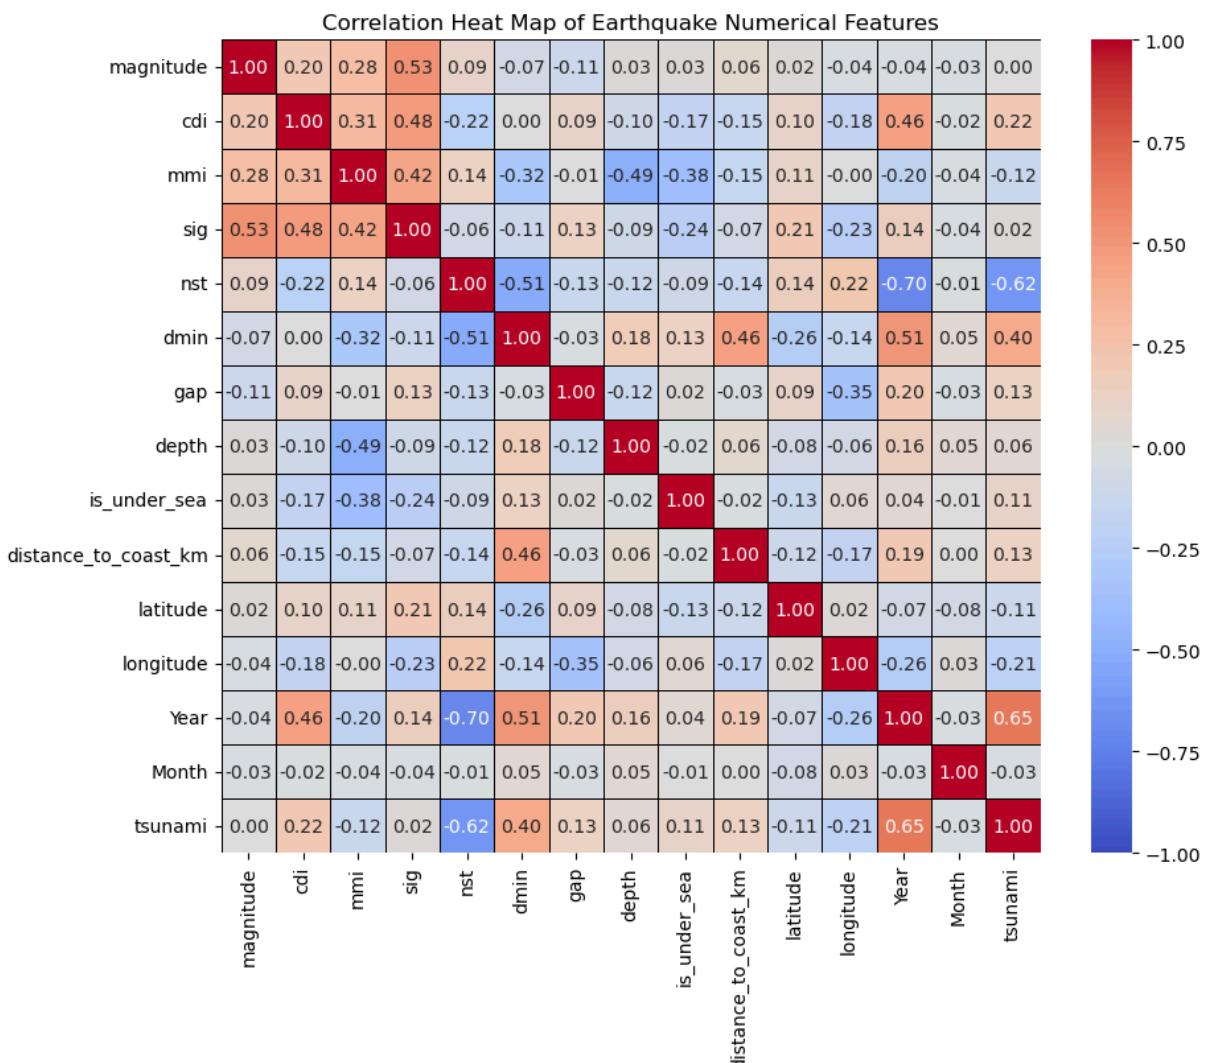
```

        'latitude', 'longitude', 'Year', 'Month', 'tsunami',
    ]
corr_df = train_df[numerical_features]

# --- 3. Calculate Correlation Matrix ---
# Calculate the correlation matrix for the selected columns
correlation_matrix = corr_df.corr()

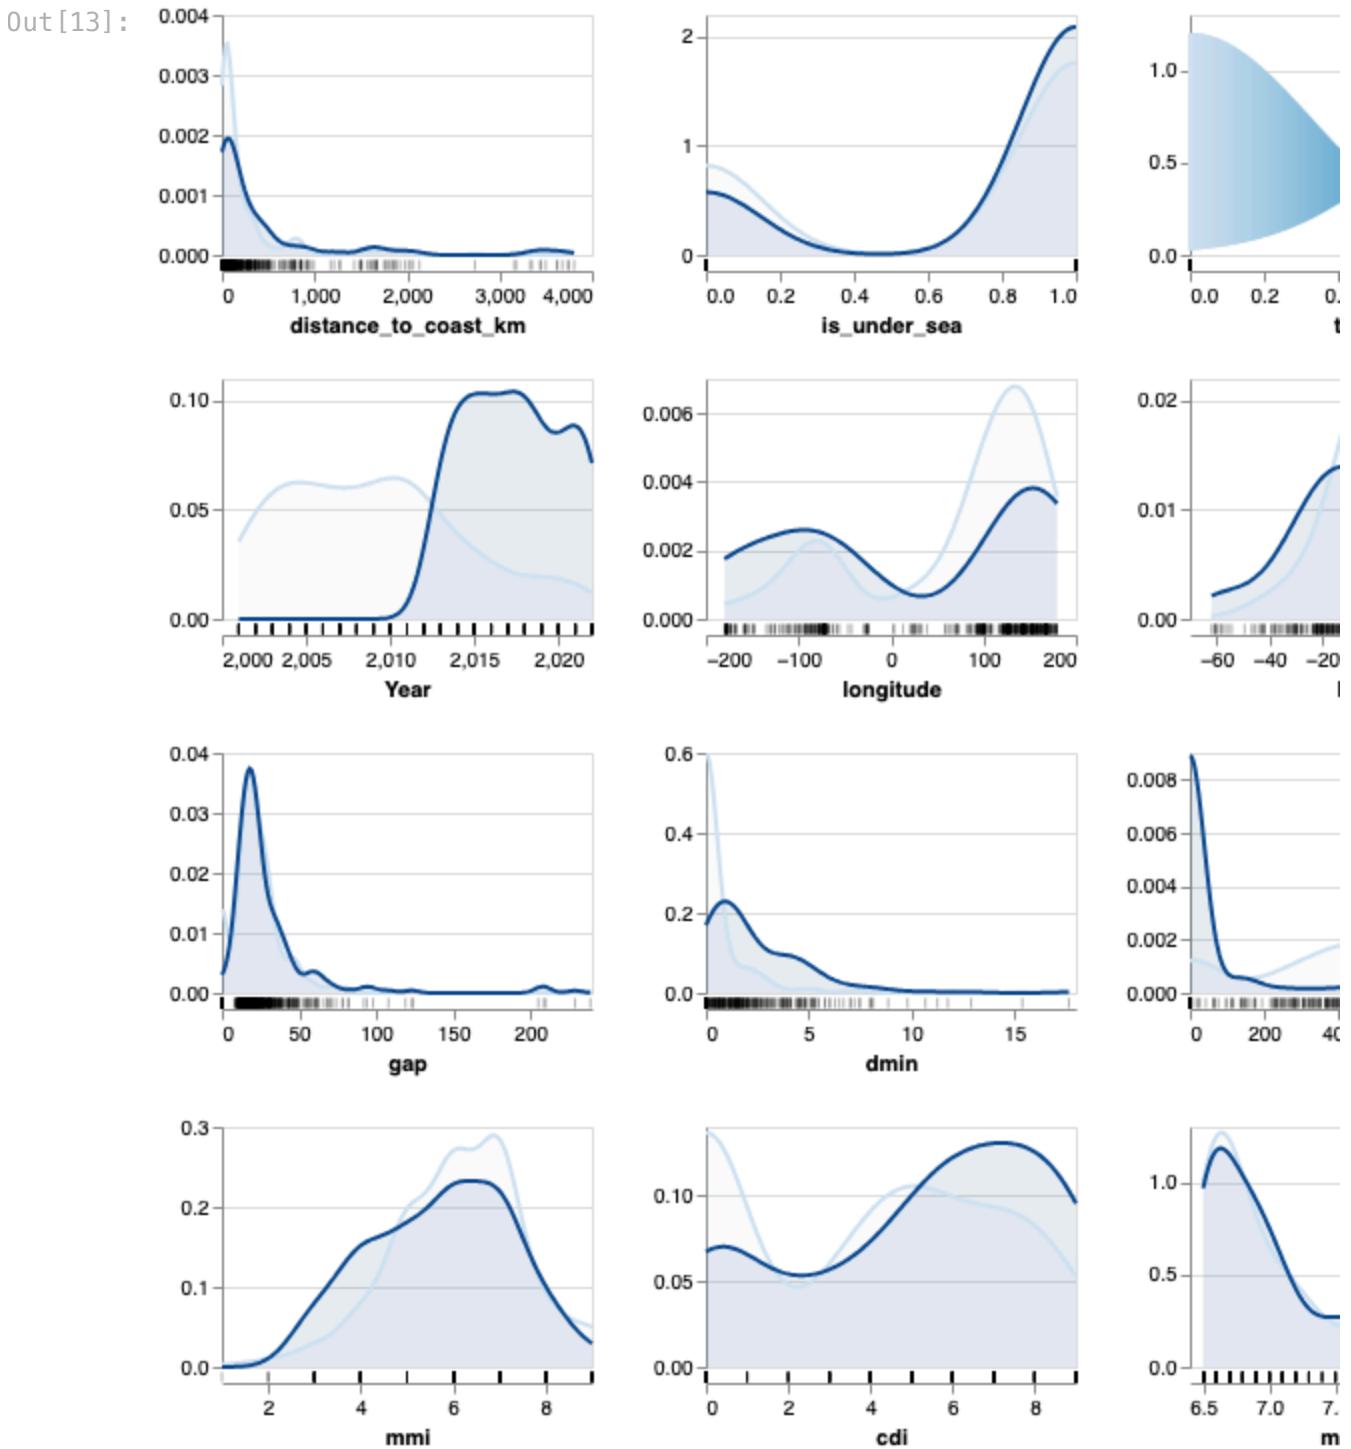
# --- 4. Generate Heat Map ---
plt.figure(figsize=(10, 8))
sns.heatmap(
    correlation_matrix,
    annot=True,           # Display correlation values
    cmap='coolwarm',      # Use a diverging color map
    fmt=".2f",            # Format numbers to two decimal places
    linewidths=0.5,       # Add lines between cells
    vmin=-1,
    vmax=1,
    linecolor='black'
)
plt.title('Correlation Heat Map of Earthquake Numerical Features')
plt.show() #

```



```
In [13]: aly.alt.data_transformers.enable('vegafusion')

aly.dist(train_df, color='tsunami', mark='area')
```



Key takeaways from EDA prior to modeling:

- It seems like there is a lack of records in tsunami before Year 2013. Since this is likely to be due to no data being recorded before 2013, we propose to get rid of year in the model. Year appears correlated simply due to missing tsunami data pre-2013, not causal effect.

- Tsunami is not seasonal, so Month should never be included in the model. It is safe to drop it.
- Numerical features which have higher correlation values with tsunami target such as cdi, dmin, , gap, and distance\_to\_coast seem to be have predictive power in classifying whether a tsunami happens or not. For others, they should still be included in the model. When they are combined with other features, they might have predictive power too.
- All the numerical features will go through StandardScaler transformation while the categorical features such as is\_under\_sea and PlateName will be transformed using OneHotEncoder.

## Machine Learning Model Selection

```
In [14]: #Function pulled from Lab 2 of DSCI571 MDS UBC, to easily extract and retrac
def mean_std_cross_val_scores(model, X_train, y_train, **kwargs):
    """
    Returns mean and std of cross validation

    Parameters
    -----
    model :
        scikit-learn model
    X_train : numpy array or pandas DataFrame
        X in the training data
    y_train :
        y in the training data

    Returns
    -----
        pandas Series with mean scores from cross_validation
    """

    scores = cross_validate(model, X_train, y_train, **kwargs)

    mean_scores = pd.DataFrame(scores).mean()
    std_scores = pd.DataFrame(scores).std()
    out_col = []

    for i in range(len(mean_scores)):
        out_col.append((f"%0.3f (+/- %0.3f)" % (mean_scores.iloc[i], std_scores
                                                    .iloc[i])))

    return pd.Series(data=out_col, index=mean_scores.index)
```

```
In [15]: X_train = train_df.drop(["tsunami"], axis=1)
y_train = train_df['tsunami']

col_transformer = make_column_transformer(
    (StandardScaler(), ['magnitude', 'depth', 'distance_to_coast_km', 'cdi',
    (OneHotEncoder(handle_unknown="ignore", sparse_output=False), ['is_under
    ("drop", ['latitude', 'longitude', 'Month', 'Year'])
```

```

        )

models = {
    "dummy": DummyClassifier(random_state=123),
    "Decision Tree": DecisionTreeClassifier(random_state=123),
    "KNN": KNeighborsClassifier(),
    "RBF SVM": SVC(random_state=123),
    "Naive Bayes": GaussianNB(),
    "Logistic Regression": LogisticRegression(max_iter=2000, random_state=123),
    "Random Forest Classifier": RandomForestClassifier(max_depth=10, random_
}

results_dict = {}

for model in models:
    pipeline = make_pipeline(col_transformer, models[model])

    results_dict[model] = mean_std_cross_val_scores(
        pipeline,
        X_train,
        y_train,
        cv=5,
        return_train_score=True)

results_df = pd.DataFrame(results_dict).T
results_df

```

Out[15]:

		fit_time	score_time	test_score	train_score
	<b>dummy</b>	0.002 (+/- 0.000)	0.001 (+/- 0.000)	0.611 (+/- 0.004)	0.611 (+/- 0.001)
	<b>Decision Tree</b>	0.004 (+/- 0.001)	0.001 (+/- 0.000)	0.813 (+/- 0.057)	1.000 (+/- 0.000)
	<b>KNN</b>	0.002 (+/- 0.000)	0.025 (+/- 0.052)	0.837 (+/- 0.036)	0.889 (+/- 0.012)
	<b>RBF SVM</b>	0.005 (+/- 0.001)	0.003 (+/- 0.000)	0.843 (+/- 0.041)	0.892 (+/- 0.009)
	<b>Naive Bayes</b>	0.002 (+/- 0.000)	0.001 (+/- 0.000)	0.534 (+/- 0.029)	0.549 (+/- 0.014)
	<b>Logistic Regression</b>	0.006 (+/- 0.004)	0.001 (+/- 0.000)	0.857 (+/- 0.040)	0.880 (+/- 0.011)
	<b>Random Forest Classifier</b>	0.054 (+/- 0.001)	0.003 (+/- 0.000)	0.870 (+/- 0.040)	0.967 (+/- 0.007)

Based on the scores obtained from 7 different models, the 3 best performing models are RBF SVM, Logistic Regression, and Random Forest Classifier, so these 3 models will be used for hyperparameter optimization. The model that achieves the highest score after hyperparameterization will be chosen to predict tsunami.

# Hyperparameter Optimization

```
In [16]: pipeline_lr = make_pipeline(col_transformer,
    LogisticRegression(max_iter=5000, random_state=123)
)

pipeline_svm = make_pipeline(col_transformer,
    SVC(kernel="rbf", probability=True, random_state=123)
)

pipeline_rf = make_pipeline(col_transformer,
    RandomForestClassifier(random_state=123)
)
```

```
In [17]: param_dist_lr = {
    "logisticregression_C": loguniform(1e-3, 1e3),
    "logisticregression_solver": ["lbfgs", "saga"],
}

param_dist_svm = {
    "svc_C": loguniform(1e-3, 1e3),
    "svc_gamma": loguniform(1e-4, 1e-1),
}

param_dist_rf = {
    "randomforestclassifier_n_estimators": range(10,50),
    "randomforestclassifier_max_depth": range(1,20),
}
```

```
In [18]: rs_lr = RandomizedSearchCV(
    estimator=pipeline_lr,
    param_distributions=param_dist_lr,
    n_iter=30,
    cv=5,
    random_state=123,
    n_jobs=-1
)

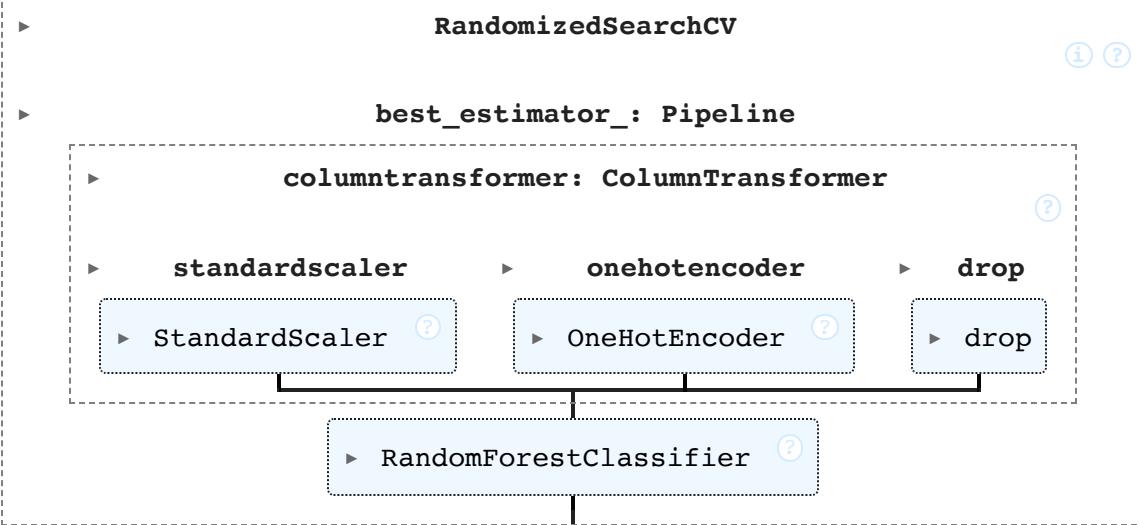
rs_svm = RandomizedSearchCV(
    estimator=pipeline_svm,
    param_distributions=param_dist_svm,
    n_iter=30,
    cv=5,
    random_state=123,
    n_jobs=-1
)

rs_rf = RandomizedSearchCV(
    estimator=pipeline_rf,
    param_distributions=param_dist_rf,
    n_iter=30,
    cv=5,
    random_state=123,
```

```
n_jobs=1  
)
```

```
In [19]: rs_lr.fit(X_train, y_train)  
rs_svm.fit(X_train, y_train)  
rs_rf.fit(X_train, y_train)
```

Out[19]:



```
In [20]: tuning_results = pd.DataFrame({  
    "Model": ["Logistic Regression", "RBF SVM", "RandomForest"],  
    "Best ROC-AUC": [rs_lr.best_score_, rs_svm.best_score_, rs_rf.best_score_],  
    "Best Params": [rs_lr.best_params_, rs_svm.best_params_, rs_rf.best_params_]})  
  
tuning_results
```

Out[20]:

	Model	Best ROC-AUC	Best Params
0	Logistic Regression	0.857314	{'logisticregression__C': 1.0256992909752545, ...}
1	RBF SVM	0.868307	{'svc__C': 178.95385188099533, 'svc__gamma': 0...}
2	RandomForest	0.872027	{'randomforestclassifier__n_estimators': 43, '...}

As seen from the table above, RandomForest seems to be performing the best compared to other models. Therefore, RandomForest with the best parameters will be chosen as the best fit model for tsunami prediction.

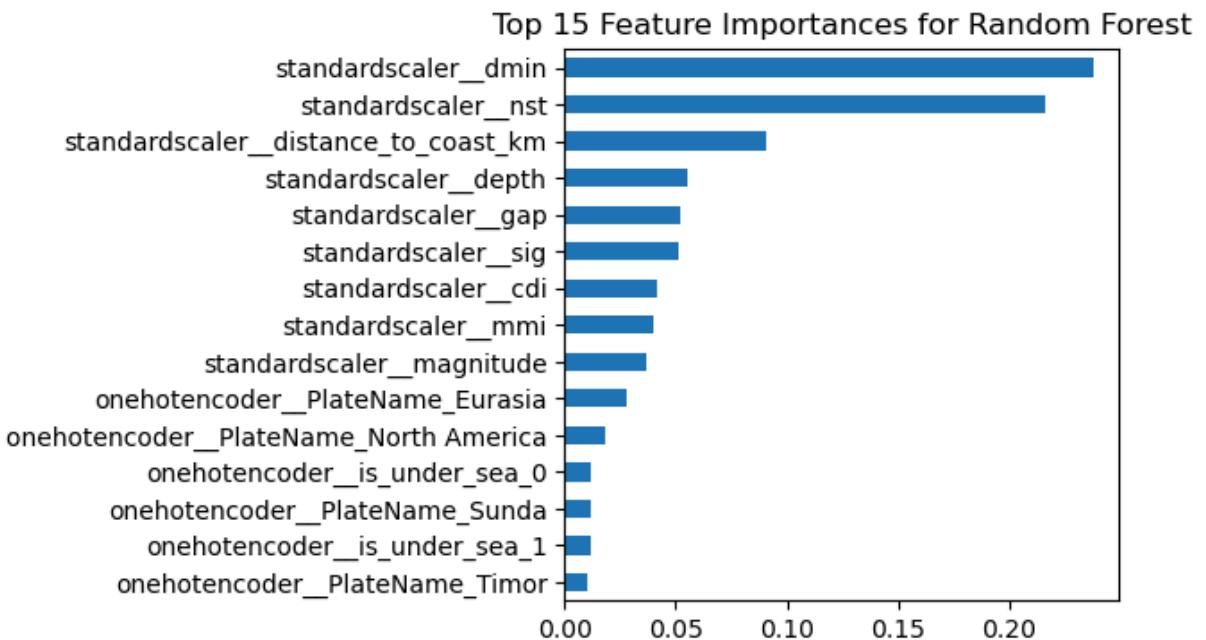
Random Forest seems to perform better compared to other models. Generally, tsunami is a natural phenomenon which is non-linear and likely to be triggered by a threshold-based process influenced by geophysical processes such as distance from coast, under-the-sea earthquakes, certain tectonic plates, etc. The engineered features in this project (distance-to-coast, under-sea classification, plate membership) create complex interactions that linear models and simple algorithms cannot capture well. Random

Forest is known for its robustness to non-linear relationships and its ability to handle both numerical and categorical inputs.

```
In [21]: # Plot top 15 feature importances for Random Forest
# Get Feature names
fitted_ct = rs_rf.best_estimator_.named_steps["columntransformer"]
feature_names = fitted_ct.get_feature_names_out()

# Get Feature importances
rf_model = rs_rf.best_estimator_.named_steps["randomforestclassifier"]
importances = pd.Series(rf_model.feature_importances_, index=feature_names)

# Plot Feature importances
importances.sort_values().tail(15).plot(kind='barh', figsize=(4, 4))
plt.title("Top 15 Feature Importances for Random Forest")
plt.show()
```



## Test Set Data Transformation

```
In [22]: # Load Natural Earth low-res land polygons via geodatasets
world_path = geodatasets.get_path("naturalearth.land")
land = gpd.read_file(world_path)

# Convert your lat/lon into point geometry
test_df['geometry'] = test_df.apply(lambda row: Point(row['longitude'], row['latitude']), axis=1)

# Make GeoDataFrame
gdf = gpd.GeoDataFrame(test_df, geometry='geometry', crs="EPSG:4326")

# Spatial join: which points intersect land?
joined = gpd.sjoin(gdf, land, how='left', predicate='intersects')

# If land data missing → ocean
```

```

test_df['is_under_sea'] = joined['featurecla'].isna().astype(int)

# Drop geometry if not needed
test_df = test_df.drop(columns=['geometry'])
test_df["distance_to_coast_km"] = test_df.apply(compute_distance_to_coast, axis=1)
test_df

```

Out[22]:

	<b>id</b>	<b>magnitude</b>	<b>cdi</b>	<b>mmi</b>	<b>sig</b>	<b>nst</b>	<b>dmin</b>	<b>gap</b>	<b>depth</b>	<b>latitude</b>	<b>longitude</b>
<b>0</b>	0	6.9	2	3	733	0	4.051	9.0	596.400	-19.7819	-178.2440
<b>1</b>	1	6.5	7	8	1297	0	2.360	15.0	29.000	-22.6784	25.1550
<b>2</b>	2	7.2	7	7	803	494	0.000	20.4	120.000	-15.5950	167.6800
<b>3</b>	3	6.5	6	6	654	0	1.640	21.0	36.000	-10.3506	161.3350
<b>4</b>	4	6.6	0	2	670	131	4.998	27.0	624.464	-25.5948	178.2780
...	...	...	...	...	...	...	...	...	...	...	...
<b>230</b>	230	6.9	7	5	747	0	1.205	21.0	35.000	0.5126	126.1890
<b>231</b>	231	7.3	8	7	921	0	3.713	19.0	165.490	-7.5924	127.5810
<b>232</b>	232	6.8	5	6	717	0	0.791	28.0	15.460	-1.8146	122.5800
<b>233</b>	233	6.9	3	4	735	0	2.785	18.0	115.000	-31.7447	-179.3730
<b>234</b>	234	7.2	5	6	802	385	0.000	27.9	39.200	6.9100	92.9580

235 rows × 15 columns

In [23]:

```

# Load Natural Earth land polygons
land_path = geodatasets.get_path("naturalearth.land")
land = gpd.read_file(land_path)

# Convert to coastline by taking the boundary
coast = land.boundary

# Project to world mercator (meters)
coast = coast.to_crs(epsg=3395)

test_df["distance_to_coast_km"] = test_df.apply(compute_distance_to_coast, axis=1)

```

In [24]:

```

gdf_test = gpd.GeoDataFrame(
    test_df,
    geometry=gpd.points_from_xy(test_df["longitude"], test_df["latitude"]),
    crs="EPSG:4326"
)

test_df = gpd.sjoin(gdf_test, plates, predicate="within", how="left")
test_df = test_df.drop(columns=['geometry', 'index_right', 'LAYER', 'Code'])
test_df

```

Out[24]:

	<b>id</b>	<b>magnitude</b>	<b>cdi</b>	<b>mmi</b>	<b>sig</b>	<b>nst</b>	<b>dmin</b>	<b>gap</b>	<b>depth</b>	<b>latitude</b>	<b>longitude</b>
<b>0</b>	0	6.9	2	3	733	0	4.051	9.0	596.400	-19.7819	-178.2440
<b>1</b>	1	6.5	7	8	1297	0	2.360	15.0	29.000	-22.6784	25.1554
<b>2</b>	2	7.2	7	7	803	494	0.000	20.4	120.000	-15.5950	167.6800
<b>3</b>	3	6.5	6	6	654	0	1.640	21.0	36.000	-10.3506	161.3350
<b>4</b>	4	6.6	0	2	670	131	4.998	27.0	624.464	-25.5948	178.2780
...	...	...	...	...	...	...	...	...	...	...	...
<b>230</b>	230	6.9	7	5	747	0	1.205	21.0	35.000	0.5126	126.1890
<b>231</b>	231	7.3	8	7	921	0	3.713	19.0	165.490	-7.5924	127.5810
<b>232</b>	232	6.8	5	6	717	0	0.791	28.0	15.460	-1.8146	122.5800
<b>233</b>	233	6.9	3	4	735	0	2.785	18.0	115.000	-31.7447	-179.3730
<b>234</b>	234	7.2	5	6	802	385	0.000	27.9	39.200	6.9100	92.9580

235 rows × 16 columns

## Model Evaluation

In [25]:

```
# 1. Use the tuned SVM model
best_model = rs_rf.best_estimator_

# 2. Prepare test features
X_test = test_df.drop(["id"], axis=1)

# 3. Predict probabilities
test_pred_proba = best_model.predict_proba(X_test)[:, 1]

# 4. Predict classes
test_pred_class = best_model.predict(X_test)

# 5. Make submission
submission = pd.DataFrame({
    "id": test_df["id"],
    "tsunami": test_pred_class
})

# 6. Save submission
submission.to_csv("submission_Nov23.csv", index=False)
```

In [26]:

```
merged_test_df = pd.merge(test_df, submission, on='id', how='inner')
# Visualize the results
plot_map_numeric(merged_test_df, "tsunami", lon_col="longitude", lat_col="latitude",
                 cmap="bwr", markersize=10, title="Tsunami Prediction")
```

### Tsunami Prediction

